

Transfomers

What happens if I learn?	What happens if I don't learn?
Expand my understanding in NLP.	I potentially lose my credibility as a Deep Learning practitioner.
Build my profile better in the NLP field.	I fall behind in the race where new tech is taking over the space with rapid innovation.
Understand the state of the art NLP language models.	Won't be able to challenge myself.
Build new products/projects with this knowledge.	Will stay behind in the race with people like AK and SP have gained a huge upper hand.
Get a better job opportunity.	Continue losing hackathons

Glossary:

Language Modeling:

Vectorizing your code

Deep Learning with Python, 2nd edition (11.4)

The Transformer Archi:

In simple terms transformer architecture leverages **neural attention** which is devoid of any recurrent or convolutional layers.

11.4.1 Understanding self-attention

When you read something/want to learn something first thing you do is skim through the content to understand the relevant bits. You devote attention to some parts depending on your interests/goals.

- Think about the **Max Pooling layer** which pools features. That's an “all or nothing” form of attention keep the most important feature and discard the rest.
- **TF-IDF normalization** assigns importance scores to tokens based on how much information different tokens are likely to carry. Important tokens get boosted while irrelevant tokens get faded out. That's a continuous form of attention.

Humans also follow the attention mechanism in some sense, when we look at an image we don't give attention to every detail present in the image. We look for some familiar features and pay attention to those details. For example: When you look at a picture of a cat, you may tend to look at the pointy ears, whiskers, the nose and eyes which enable us to classify the picture as a cat.

Illustration in the image below:

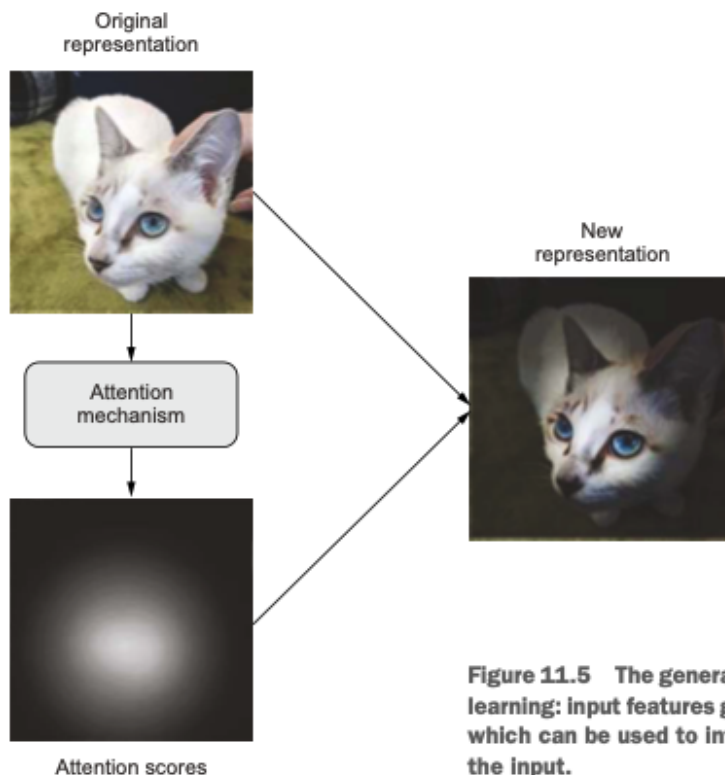


Figure 11.5 The general concept of “attention” in deep learning: input features get assigned “attention scores,” which can be used to inform the next representation of the input.

Then an interesting question is,

How do you determine which features needs the most/ least attention?

As a simple solution you can score your features with relevant features getting higher scores and the less relevant features being assigned lesser scores.

Now this leads to the next questions:

How these scores should be computed? and

What you should do with these scores ?

Ans: It will vary from approach to approach. The crux of the idea is to make features **content-aware**.

Vector spaces that capture the “shape” of the semantic relationships between different words.

In an embedding space, a single word has a fixed position-a fixed set of relationships with every other word in the space.

But language doesn't work like that, the meaning of a word is usually context-specific.

Ex:- *When you mark the date, you're not talking about the same "date" as when you go on a date, nor is it the kind of date you'd buy at the market. When you say, "I'll see you soon," the meaning of the word "see" is subtly different from the "see" in "I'll see this project to its end" or "I see what you mean." And, of course, the meaning of pronouns like "he," "it," "in," etc., is entirely sentence-specific and can even change multiple times within a single sentence.*

Therefore we need vector representation of words depending on the words that surround it. That's where self-attention comes in.

The purpose of **self-attention** is to modulate the representation of a token by using the representations of related tokens in the sequence. This produces **context aware token representations**.

Consider an example sentence: "**The train left the station on time.**" Now, consider one word in the sentence: station. What kind of station are we talking about? Could it be a radio station? Maybe the International Space Station? Let's figure it out algorithmically via self-attention (see figure 11.6).

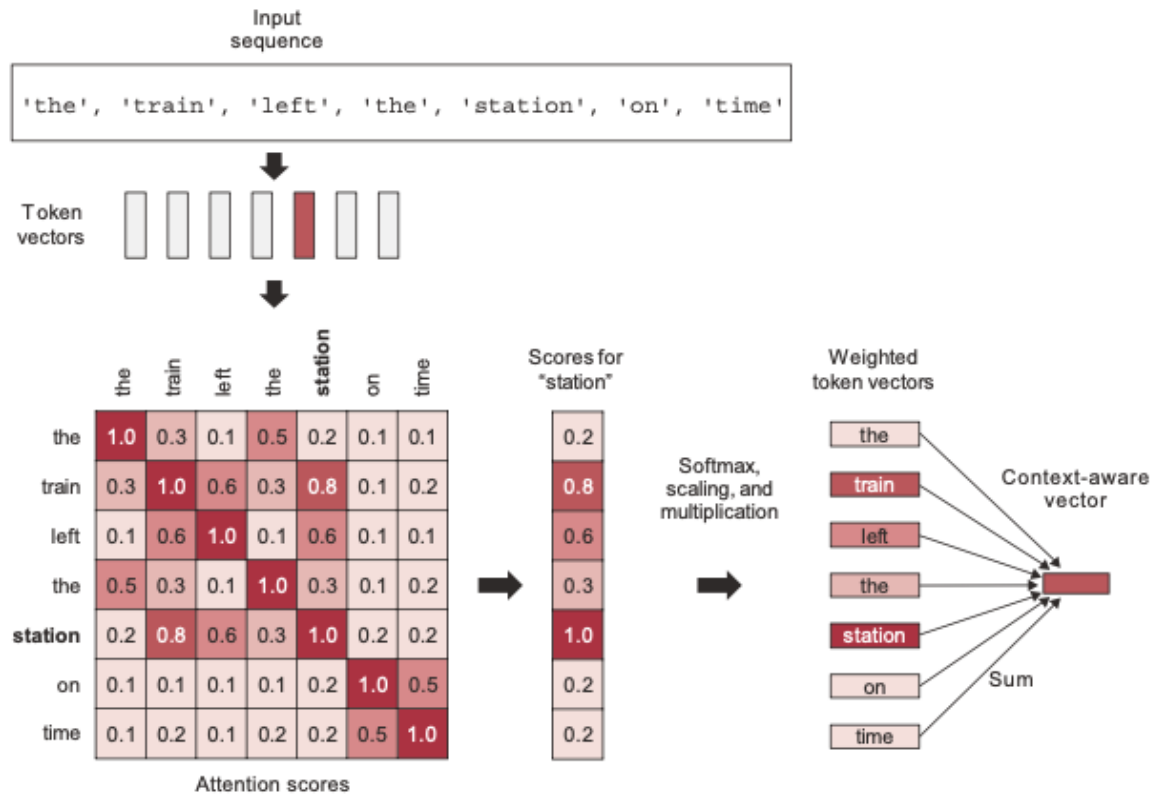


Figure 11.6 Self-attention: attention scores are computed between “station” and every other word in the sequence, and they are then used to weight a sum of word vectors that becomes the new “station” vector.

Algorithm Steps:

1. Compute relevancy score (cosine -used because computationally very efficient distance function-) between the vector for “**station**” and every other vector in the sentence. These scores are our **attention scores**. Later these scores will go through a scaling function and a softmax.
2. Then compute the sum of all word vectors in the sentence, weighted by the relevancy score. Words closely related to “station” will contribute more to the sum. Then the resulting vector is our new representation for the word “station”.

Then repeat this process for every word in the sentence producing a new sequence of vectors encoding the sentence.

At this point I have a few questions:

1. How are the initial vector embeddings calculated ?

2. When does the vector calculation stop ? After reading each sentence or the whole corpus?

Please read this before resuming again: <https://towardsdatascience.com/illustrated-self-attention-2d627e33b20a>

Keras has a built-in layer to handle attention: the **MultiHeadAttention** layer. Here's how you would use it:

```
num_heads = 4
embed_dim = 256
mha_layer = MultiHeadAttention(num_heads=num_heads, key_dim=embed_dim)
outputs = mha_layer(inputs, inputs, inputs)
```

Reading this, you're probably wondering

- Why are we passing the inputs to the layer *three* times? That seems redundant.
- What are these “multiple heads” we're referring to? That sounds intimidating do they also grow back if you cut them?

Both of these questions have simple answers. Let's take a look.